

# Today's tutorial

- How to make inference/training fast?
- How to deal with limited hardware?
- How to get more compute for bigger projects?

# How to make inference fast?

- Flash Attention
- Quantization
- BetterTransformer
- Optimum library

# Flash Attention

- Time and memory complexity of self-attention is quadratic w.r.t. sequence length
  - The longer the sequence, the more time inference/training takes
- Flash attention introduces additional parallelization over the sequence length to mitigate this
- Smarter partitioning of the workload between GPU threads -> less reads and writes from memory
- Exact algorithm, extension to block-sparse is non-exact
- Only usable for dtypes fp16 and bf16 and on NVIDIA GPUs
- Experimental, only usable without padding tokens

```
pip install flash-attn --no-build-isolation
```

To enable FlashAttention-2, add the `use_flash_attention_2` parameter to `from_pretrained()`:

```
import torch
from transformers import AutoModelForCausalLM, AutoTokenizer, LlamaForCausalLM

model_id = "tiiuae/falcon-7b"
tokenizer = AutoTokenizer.from_pretrained(model_id)

model = AutoModelForCausalLM.from_pretrained(
    model_id,
    torch_dtype=torch.bfloat16,
    use_flash_attention_2=True,
)
```

# Quantization

- Computers are bad at floating point arithmetic
- We need to choose a certain level of precision whenever we deal with floating point numbers
- Neural networks are just large collections of floating point numbers
- If we reduce the precision of the model, we need much less memory and speed up matrix multiplication considerably
- We may lose some accuracy → applicability dependent on use case

# Quantization



Downloads last month  
**2,886,056**



 **Safetensors** ⓘ

Model size

407M params

Tensor type

F32



# Quantization: Common Data Types

Data Type	Accumulates to
float16	float16
bfloat16	float32
int16	int32
int8	int32

# Accumulation Data Types

- Suppose we want to add two int8 values,  $A=127$ ,  $B=127$ :
- $A + B = C$
- 127 is the maximum value representable in int8 → accumulate to bigger data type to prevent large precision losses



# Quantization with float16

- Same representation as float32 → straightforward quantization
- Questions to answer beforehand:
  - Do operations have a float16 implementation?
  - Does the hardware support operations in float16?
  - Is my operation sensitive to lower precision?
- Some operations in ML are sensitive to changes in very low values (e.g. layer norm)
- With larger scale generative LLMs, quantization to float16 is pretty standard

# Quantization to int8

- 256 values representable in int8
- Find best way to project range of float32 value to int8 space
- Consider float  $x$  in range  $[a, b]$ , using affine quantization scheme:

$$X = s^*(x_q - Z)$$

- $x_q$  = quantized int8 value associated to  $x$
- $S$  = defined as  $(\text{float\_max} - \text{float\_min}) / (\text{int\_max} - \text{int\_min})$
- $Z$  = zeropoint, int8 value corresponding to 0 in float32 realm
  - We do not want to make errors when mapping zero!

# Quantization to int8

- Compute  $x_q$ :

$$x_q = \text{round}(x/S + Z)$$

- All values outside the  $[a,b]$  range are clipped to the closest representable value:

$$x_q = \text{clip}(\text{round}(x/S + Z), \text{round}(a/S + Z), \text{round}(b/S + Z))$$

# Quantization in Huggingface Transformers

```
# these versions support 8-bit and 4-bit  
pip install bitsandbytes>=0.39.0 accelerate>=0.20.0
```

```
from transformers import AutoModelForCausalLM  
  
model_name = "bigscience/bloom-2b5"  
model_4bit = AutoModelForCausalLM.from_pretrained(model_name, device_map="auto", load_in_4bit=True)
```

```
from transformers import AutoModelForCausalLM  
  
model_name = "bigscience/bloom-2b5"  
model_8bit = AutoModelForCausalLM.from_pretrained(model_name, device_map="auto", load_in_8bit=True)
```

# BetterTransformer

- Optimized execution of Huggingface Transformer functions:
  - Fusing multiple sequential operations into single “kernel” to reduce number of computations performed
  - More performant handling of padding tokens

```
python -m pip install optimum
```

```
model = model.to_bettertransformer()
```

```
model = model.reverse_bettertransformer()  
model.save_pretrained("saved_model")
```

# Optimum

- Similar optimizations to BetterTransformer
- Only for nvidia GPUs

```
from optimum.onnxruntime import ORTModelForSequenceClassification

ort_model = ORTModelForSequenceClassification.from_pretrained(
    "distilbert-base-uncased-finetuned-sst-2-english",
    export=True,
    provider="CUDAExecutionProvider",
)
```

# How to make training fast?

- Either speed up computation, optimize memory utilization or both

Method/tool	Improves training speed	Optimizes memory utilization
<u>Batch size choice</u>	Yes	Yes
<u>Gradient accumulation</u>	No	Yes
<u>Gradient checkpointing</u>	No	Yes
<u>Mixed precision training</u>	Yes	(No)
<u>Optimizer choice</u>	Yes	Yes
<u>Data preloading</u>	Yes	No
<u>DeepSpeed Zero</u>	No	Yes
<u>torch.compile</u>	Yes	No

# Batch Size Choice

- Always the starting point
- Batch size too large: OOM error
- Batch size too small: inefficient memory usage, longer training time, more cost
- Has to be of size  $2^N$
- Smaller batch sizes usually converge faster than large batch sizes



# Gradient Accumulation

- Computes gradient at smaller steps than for the entire batch at once
- Iteratively compute gradients by performing forward/backward passes through the model
- Accumulate gradients during process
- Once enough gradients have been accumulated, perform model's optimization
- Enables larger batch size than would usually be possible
- But: training process becomes slower → balance gradient accumulation steps and batch size

# Gradient Accumulation and Batch Size in Transformers

```
training_args = TrainingArguments(per_device_train_batch_size=1, gradient_accumulation_steps=4,
```

# Gradient Checkpointing

- Saving all activations from forward pass result in large memory overhead
- Disregarding and recomputing activations during the backward pass results in large computational overhead
- Gradient Checkpointing strategically select the activations to keep, resulting in only a fraction of the activations to be recomputed

```
training_args = TrainingArguments(  
    per_device_train_batch_size=1, gradient_accumulation_steps=4, gradient_checkpointing=True,  
)
```

# Mixed Precision Training

- Reduces precision of certain variables to achieve computational speed up
- Decreases training time, but can also increase GPU memory utilization, because two representations of the model have to be saved
- Data type to be used dependent on hardware support

```
training_args = TrainingArguments(per_device_train_batch_size=4, fp16=True,
```

```
training_args = TrainingArguments(bf16=True, **default_args)
```

# Optimizer Choice

- Optimizes a function (our neural network) w.r.t. an objective function (our loss function)
- To do this, the optimizer has to make changes to all of the weights and biases
- Large impact on training performance/convergence

# Optimizer Choice

- Default is Adam(W) which stores rolling averages of the previous gradients
- This results in significant increase of memory usage
- Out of the box alternatives for huggingface transformers:

Optimizer	GPU Memory for 3B Model
AdamW	24GB
Adafactor	12GB
8bit BNB quantized	6GB

# Adafactor

- Aggregates the rolling averages Adam uses by summing rolling average row and column wise
- Possibly slower convergence
- However, significant reductions in memory throughput can be achieved

```
training_args = TrainingArguments(per_device_train_batch_size=4, optim="adafactor",
```

# 8-bit Adam

- Quantizes optimizer states instead of aggregating them

```
training_args = TrainingArguments(per_device_train_batch_size=4, optim="adamw_bnb_8bit",
```



# Sharding

- Usable for inference and training
- Divide layers of neural network into single files
- Load files one by one on the GPU during training or inference
- Also used for large distributed LLM applications
- Multiple ways of doing this with huggingface

# Two Cases

- I have one GPU and its too small
- I have multiple GPUs and want to distribute inference/training
  - [https://huggingface.co/docs/accelerate/concept\\_guides/big\\_model\\_inference](https://huggingface.co/docs/accelerate/concept_guides/big_model_inference)

# Other techniques



- Data preloading: makes sure the GPU loads as much data as it can possibly handle
  - Set `dataloader_num_workers` in `TrainingArguments` to higher value if GPU utilization is far from 100%
- `Torch.compile`: uses some low-level computational graph magic to optimize training

```
training_args = TrainingArguments(torch_compile=True, **default_args)
```

# General Advice

- Working with LLMs, especially training, is a applying lot of trial and error while trying to keep your nerves
- Despite all the fancy optimization techniques, more compute is almost always better
- Fine-tuned models can match much larger models on their respective task in a lot of cases
- Guardrails can make models significantly worse

# Censored models getting outperformed?



# General Advice

- Cloud Providers may have some starting credits to get you hooked (\$300 on GCP → great for bigger projects, thesis)
- Being able to access servers remotely and letting your stuff run without having to keep a local machine on is a major advantage

# Compute for Assignment 2

- We will rent a 32GB RAM Machine for you to use
- You will have your own user accounts and SSH access
- You will be able to run python scripts/install python packages
- This is (kind of) an experiment, please do not crash the server, distribute your SSH keys anywhere, etc.
- Windows users may install WSL2 to log onto the machine